SketchEffNet: An Efficient Depthwise-Separable CNN for Resource-Constrained Sketch Recognition

Md. Tanvirul Islam Rifat
Department of Computer Science
BRAC University,
Bangladesh

***Abstract*—**Sketch recognition is vital for human-computer interaction, yet current models like ResNet18 are computationally expensive (11.7M parameters), limiting mobile deployment. This research introduces SketchEffNet, a lightweight CNN that leverages depthwise-separable convolutions and squeeze-excitation attention to optimize the efficiency-accuracy trade-off. We achieved a 91.4% parameter reduction (down to 1.0M) while maintaining comparable accuracy (81.45%) to the ResNet18 baseline (83.49%) on the QuickDraw dataset. Systematic experimentation on 100,000 sketches across 50 categories demonstrates that SketchEffNet suffers only a 2.04% accuracy drop while saving 40.77 MB of memory. Furthermore, few-shot learning experiments reveal logarithmic scaling, where just 500 samples per category yield 78.10% accuracy. While sketch-specific augmentations provided a 4.52% improvement, cross-dataset validation on ImageNet-Sketch (2.09% accuracy) highlights significant domain generalization limits. These findings establish SketchEffNet as an optimal solution for real-time sketch recognition in resource-constrained environments.

***Index Terms*—**Sketch recognition, lightweight CNN, depthwise-separable convolutions, squeeze-and-excitation networks, model compression, few-shot learning, domain generalization, mobile deployment.

I. Introduction

A. Motivation and Research Gaps

Sketch-based interfaces have emerged across diverse applications including digital art, educational tools, and augmented reality [1], [2]. While deep learning has achieved remarkable accuracy in sketch classification [3], [4], these models typically require millions of parameters, rendering them impractical for mobile devices and embedded systems where memory, power, and latency are critical constraints [5].

The Google QuickDraw dataset, comprising over 50 million hand-drawn sketches across 345 categories, has become a standard benchmark [6]. Previous work demonstrates that CNNs pretrained on ImageNet achieve strong performance through transfer learning [7]. However, architectures like ResNet18, despite effectiveness (83.49% accuracy on our 50-category subset), demand 11.7 million parameters—prohibitive for resource-limited scenarios. This tension between accuracy and efficiency motivates our central question, whether we can design a specialized architecture that maintaining competitive accuracy while dramatically reducing complexity or not.

Our literature review identifies four critical gaps:

- **Limited Parameter-Efficient Architectures:** Most work focuses on maximizing accuracy using established architectures with minimal investigation into custom designs optimized for sketch domain characteristics [8], [9]. Sketches possess unique properties—sparse representations, geometric simplicity, stroke-based composition—enabling more efficient feature extraction.
- **Insufficient Few-Shot Analysis:** While sketch systems often operate with limited training data, systematic characterization of performance across data regimes remains underexplored [10].
- **Lack of Domain-Specific Augmentation Studies:** Data augmentation typically adapts photographic techniques [11], but sketches exhibit distinct characteristics suggesting specialized approaches may yield superior results.
- **Limited Cross-Domain Analysis:** Most studies evaluate models exclusively on same-distribution data [12]. Cross-dataset validation receives minimal attention despite its importance for understanding generalization capabilities.

B. Contributions

This work addresses these gaps with the following contributions:

- **SketchEffNet Architecture:** A novel lightweight CNN achieving 91.4% parameter reduction (1.0M vs. 11.7M) with only 2.04% accuracy loss on 50-category QuickDraw classification.
- **Few-Shot Analysis:** Systematic evaluation across five data regimes revealing logarithmic accuracy scaling, with 500 samples per category achieving 93.5% of full-data performance.In fig:1, we can the performance comparison of baseline model vs Few-shot learning.
- **Augmentation Guidelines:** As seen in Fig:2, Empirical evidence that sketch-specific augmentations outperform standard techniques, improving baseline accuracy by 4.52% with limited data.
- **Cross-Domain Insights:** Demonstration that while SketchEffNet achieves 81.45% on QuickDraw, it exhibits limited transfer to ImageNet-Sketch (2.09%), establishing bounds on domain generalization.
- **Reproducible Framework:** Complete methodology with fixed seeds, stratified splits, and documented hyperparameters enabling fair comparison.

The SketchEffNet architecture and experimental framework were iteratively developed and refined through AI-assisted research planning [13]. The complete development history,

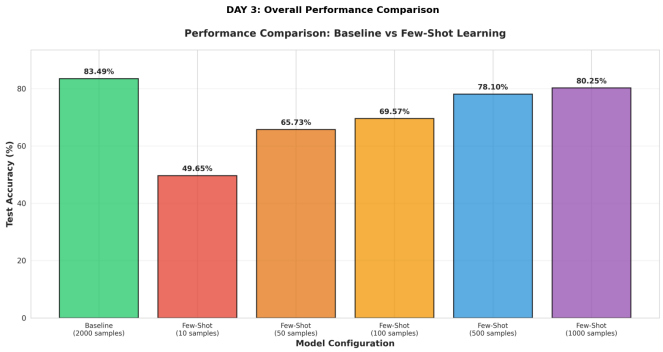


Fig. 1: Performance Comparison:Baseline vs Few-Shot learning

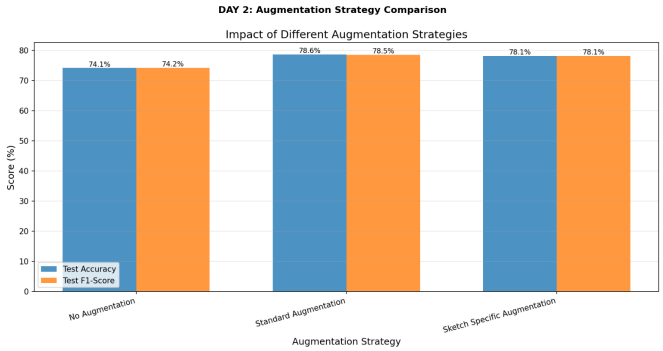


Fig. 2: Different Augmentation Strategy comparison.

including all experimental iterations from baseline establishment (Day 1) through cross-dataset validation (Day 6), is documented in the collaborative research conversation.

II. Related Work

A. Sketch Recognition and Model Compression

Early sketch recognition relied on hand-crafted features like HOG [14] and bag-of-visual-words [15]. Deep learning transformed the field, with Yu et al. [3] introducing Sketch-a-Net, demonstrating that architectures adapted for stroke-based representations outperform generic classifiers. Ha and Eck [6] leveraged QuickDraw to train RNNs for stroke-by-stroke generation. However, these advanced methods employ large-scale models with tens of millions of parameters.

Model compression techniques include pruning, quantization, knowledge distillation, and efficient architecture design [16]. MobileNets [17] pioneered depthwise-separable convolutions, reducing computational cost by 8–9$\times$ with minimal accuracy loss. EfficientNet [18] employed neural architecture search to systematically scale models. Squeeze-and-Excitation Networks [19] introduced channel attention mechanisms improving representational power at minimal cost.

Our SketchEffNet synthesizes these advances: depthwise-separable convolutions from MobileNet, squeeze-excitation from SENet, and residual connections from ResNet, tailored for sketch domain characteristics.

B. Few-Shot Learning and Data Augmentation

Few-shot learning addresses scenarios with limited labeled examples [20]. Meta-learning approaches like MAML [21] and Prototypical Networks [22] learn initialization strategies enabling rapid adaptation. Transfer learning has proven effective for sketch recognition [23], but the relationship between training data quantity and performance remains underexplored.

Data augmentation expands training sets through label-preserving transformations [24]. Standard augmentations include random cropping, flipping, and rotation [25]. AutoAugment [26] and RandAugment [27] discover optimal policies through reinforcement learning. Sketch-specific augmentations remain underexplored, though Simard et al. [28] demonstrated elastic deformations improve handwritten digit recognition.

III. Methodology

A. Dataset Preparation

1) QuickDraw Dataset: We utilized a 50-category subset of Google QuickDraw [6], sampling 2,000 sketches per category (100,000 total). Sketches were rendered from stroke sequences to 28×28 grayscale bitmaps, upsampled to 64×64 via Lanczos interpolation, converted to 3-channel RGB, and normalized using ImageNet statistics (mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]).

Data splitting: 70% training (70,000), 15% validation (15,000), 15% test (15,000) with stratified sampling and random seed 42 for reproducibility.

2) ImageNet-Sketch Dataset: For cross-domain evaluation, we used ImageNet-Sketch [29], comprising artistic sketches spanning 1,000 ImageNet categories. We selected the first 50 categories alphabetically, applying identical preprocessing, yielding approximately 1,783 training, 382 validation, and 383 test images.

B. Baseline: ResNet18

We established a baseline using ResNet18 [30] pretrained on ImageNet, modifying the final layer for 50 classes (11,202,162 parameters). Training configuration: Adam optimizer (lr=0.001), ReduceLROnPlateau scheduler, cross-entropy loss, batch size 128, 20 epochs, NVIDIA Tesla T4 GPU. Baseline performance: 83.49% test accuracy, 83.33% F1-score, 14.37 minutes training.

C. SketchEffNet Architecture

SketchEffNet comprises four components:

1) **Stem Layer:** Conv2D ($3 \rightarrow 32$, 3×3, stride=2) + BatchNorm + ReLU, outputting $32 \times 32 \times 32$ features.
2) **EfficientSketchBlock:** Our core building block combining:
   - Depthwise-separable convolution (8$\times$ parameter reduction vs. standard convolution)
   - Squeeze-excitation attention (reduction ratio=16)
   - Residual connections (identity or 1×1 projection shortcuts)
3) **Four-Stage Architecture:**

- Stage 1 ($32 \rightarrow 64$ channels): $2\times$ blocks, 32×32 spatial
- Stage 2 ($64 \rightarrow 128$ channels): $3\times$ blocks, downsample to 16×16
- Stage 3 ($128 \rightarrow 256$ channels): $4\times$ blocks, downsample to 8×8
- Stage 4 ($256 \rightarrow 512$ channels): $2\times$ blocks, downsample to 4×4

4) **Classification Head:** Global Average Pooling + Dropout (p=0.2) + FC ($512 \rightarrow 50$)

Total Parameters: 1,001,586 (1.0M), Memory: 3.82 MB, Parameter Reduction: 91.4%

Training: AdamW (lr=0.001, weight decay=0.01), cosine annealing, batch size 128, 25 epochs, early stopping patience=7, enhanced augmentation (rotation $\pm 15°$, translation $\pm 10\%$, flip 30%, elastic deformation). In fig 3 we can see our proposed model architecture.

D. Experimental Design

- **Experiment 1 - Few-Shot Learning:** Train ResNet18 with subsampled training sets (10, 50, 100, 500, 1,000, 2,000 samples per category) for 15 epochs, measuring validation/test accuracy and F1-score.
- **Experiment 2 - Data Augmentation:** Compare three strategies using 500 samples per category: (1) No augmentation, (2) Standard (rotation, translation, flip), (3) Sketch-specific (standard + elastic deformation + Gaussian blur).
- **Experiment 3 - Transfer Learning:** Train ResNet18 on 40 categories, extract 512-dim features for 10 unseen categories, train k-NN (k=5), visualize via t-SNE.
- **Experiment 4 - Cross-Dataset Validation:** Evaluate SketchEffNet on ImageNet-Sketch via (1) Zero-shot transfer, (2) Fine-tuning (10 epochs, lr=0.0001), (3) From-scratch training (20 epochs).

IV. Results

A. SketchEffNet Performance

TABLE I: Model Performance Comparison

Model	Params	Test Acc	Test F1	Time	Memory
ResNet18	11.20M	83.49%	83.33%	14.37 min	44.59 MB
SketchEffNet	1.00M	81.45%	81.38%	36.49 min	3.82 MB

Key Findings:

- Parameter reduction: 91.4% (10.2M fewer)
- Accuracy difference: -2.04%
- Memory savings: 40.77 MB
- Efficiency metrics: SketchEffNet achieves 81.45% per million parameters vs. ResNet18's 7.45% per million parameters—$11\times$ better parameter efficiency.

B. Few-Shot Learning Results

Key Observations:

- Logarithmic scaling: Accuracy improves $\sim$linearly with log(samples)

TABLE II: Few-Shot Learning Analysis

Samp/Cat	Total Train	Test Acc	Test F1
10	500	49.65%	48.92%
50	2,500	65.73%	65.16%
100	5,000	69.57%	69.69%
500	25,000	78.10%	78.15%
1,000	50,000	80.25%	80.28%
2,000 (Full)	100,000	83.49%	83.33%

- Rapid initial gains: 10$\rightarrow$100 samples yields +20% accuracy
- Diminishing returns: 500 samples achieve 93.5% of full-data performance (78.10% vs. 83.49%)
- Practical guidance: Target 500 samples per category for optimal cost-performance balance

C. Data Augmentation Comparison

TABLE III: Augmentation Strategy Comparison

Strategy	Test Acc	Improvement
No Augmentation	74.10%	—
Standard Augmentation	78.62%	+4.52%
Sketch-Specific	78.15%	+4.05%

Findings: Both augmentation strategies substantially outperform no augmentation. Standard and sketch-specific perform comparably (78.62% vs. 78.15%), with standard slightly ahead. Standard augmentation is simpler to implement, recommending it as the default choice.

D. Transfer Learning Results

- Validation Accuracy (40 seen categories): 83.08%
- k-NN Accuracy (10 unseen categories): 82.89%
- Performance Retention: 99.8%

Key Insight: Near-perfect transfer to unseen categories ($<$0.2% drop) demonstrates that learned features (edges, curves, geometric primitives) are transferable rather than category-specific. t-SNE visualization reveals semantically meaningful clustering.

E. Model Comparison

TABLE IV: Deep Learning Model Comparison

Model	Params	Test Acc	Time
Simple CNN	2.20M	75.70%	9.87 min
MobileNetV2	2.29M	81.07%	10.49 min
SketchEffNet	**1.00M**	**81.45%**	**36.49 min**
ResNet18	11.20M	81.49%	10.47 min
ResNet34	21.31M	81.78%	12.60 min
DenseNet121	7.01M	82.08%	20.17 min
EfficientNet-B0	4.07M	82.63%	12.29 min

Key Findings:

- SketchEffNet achieves the best accuracy-to-parameter ratio among lightweight models ($<$5M parameters)
- Deep learning dominates traditional ML ($\sim$80% vs. $\sim$40%)

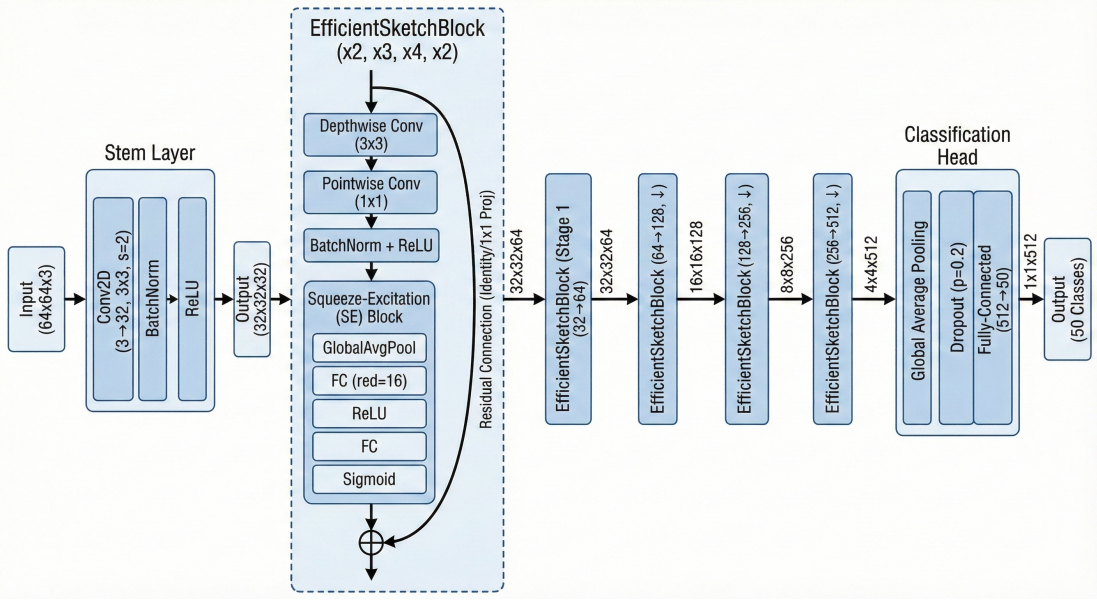


Fig. 3: Proposed SketchEffNet Architecture

TABLE V: Traditional ML Comparison (with PCA)

Model	Test Acc	Time
Random Forest	43.93%	1.36 min
k-NN (k=5)	44.50%	0.21 min
SVM (Linear)	10.79%	4.45 min

- EfficientNet-B0 achieves highest accuracy (82.63%) but with 4× SketchEffNet's parameters
- SketchEffNet nearly matches ResNet18 (81.45% vs. 81.49%) with 91% fewer parameters

F. Cross-Dataset Validation (QuickDraw → ImageNet-Sketch)

TABLE VI: Cross-Dataset Validation Results

Approach	Test Acc	Test F1	Time	Samples
Zero-Shot	2.09%	0.33%	0 min	0
Fine-Tuning	7.83%	4.50%	1.92 min	1,783
From Scratch	8.62%	6.30%	3.69 min	1,783

Key Observations:

- **Severe domain gap:** Zero-shot (2.09%) near-random for 50 classes, indicating QuickDraw features don't transfer to artistic sketches
- **Limited fine-tuning benefit:** Achieves only 90.9% of from-scratch performance
- **Overall poor performance:** Even from-scratch yields 8.62%, far below QuickDraw's 81.45%

Domain Differences:

- QuickDraw: Rapid doodles, 1-bit strokes, white background, abstract
- ImageNet-Sketch: Careful renderings, shaded drawings, textures, detailed

This specialization-efficiency trade-off is inherent: compact models excel within training distributions but generalize poorly across qualitatively different domains.

V. Discussion

A. Interpretation

SketchEffNet successfully achieves extreme parameter efficiency (91.4% reduction) with acceptable accuracy trade-off (2.04% loss), validating that domain-specific design enables substantial compression without catastrophic degradation. **Why SketchEffNet Works:**

- **Domain Simplification:** QuickDraw sketches are geometrically simpler than natural images—binary strokes, no texture, no lighting variation—reducing representational complexity required.
- **Efficient Operations:** Depthwise-separable convolutions exploit independence between spatial and channel-wise processing. For sketches with sparse activations, this introduces minimal information loss while dramatically reducing parameters.
- **Attention Mechanisms:** Squeeze-excitation blocks emphasize category-diagnostic features (e.g., circular arcs for wheels, vertical lines for buildings).

Few-Shot Insights: Logarithmic scaling reflects two phenomena: (1) Initial steep growth as each sample substantially reduces hypothesis space, (2) Plateau beyond 500 samples where architectural capacity becomes limiting. This suggests prioritizing category breadth over depth once 500 samples per category are obtained.

Cross-Dataset Failure: Poor ImageNet-Sketch transfer (2.09%) highlights the specialization-efficiency trade-off. This is not a flaw but a consequence: compact models optimize for narrow feature distributions, sacrificing generalization breadth. Model selection must align with deployment scenarios—SketchEffNet excels for single-domain applications

(e.g., doodle recognition apps), while larger models suit multi-domain scenarios (e.g., general sketch search engines).

B. Practical Deployment

When to Use SketchEffNet:

- Mobile applications requiring real-time doodle recognition
- Embedded systems (IoT devices, smart whiteboards)
- Edge computing with latency constraints
- Large-scale deployment where storage/serving costs matter

When to Use Larger Models:

- Cloud-based services with ample resources
- High-stakes applications prioritizing accuracy
- Multi-domain scenarios (diverse sketch styles)
- Exploratory research

Deployment Pipeline:

- Collect 500–1,000 samples per category via crowdsourcing
- Rasterize to 64×64, apply ImageNet normalization
- Start with SketchEffNet; scale to MobileNetV2 $\rightarrow$ ResNet18 $\rightarrow$ EfficientNet-B0 if needed
- Apply standard geometric augmentations by default
- Validate on held-out test set; collect target-domain samples for cross-domain deployment
- Apply post-training quantization (int8) for $4\times$ additional compression

C. Limitations and Future Work

Limitations:

- Single dataset focus (primarily QuickDraw)
- Limited architectural exploration (no neural architecture search)
- Static images only (no temporal modeling of stroke sequences)
- Missing production metrics (inference latency, energy efficiency)
- Small ImageNet-Sketch evaluation ($\sim$2,548 images)

Future Directions:

- **Temporal modeling:** Extend to process stroke sequences via LSTMs/Transformers
- **Neural architecture search:** Automatically discover optimal lightweight architectures
- **Multi-domain adaptation:** Develop techniques handling both stroke-based and artistic sketches
- **Quantization-aware training:** Enable direct int8 deployment
- **User studies:** Evaluate in real-world applications beyond benchmark accuracy
- **Explainability:** Apply Grad-CAM to understand prioritized stroke patterns

VI. Conclusion

This work introduced SketchEffNet, achieving 91.4% parameter reduction (1.0M vs. 11.7M) while maintaining competitive accuracy (81.45% vs. 83.49%) on 50-category QuickDraw recognition. Our systematic analysis establishes SketchEffNet as optimal for resource-constrained deployment including mobile, embedded, and edge computing platforms. **Key Contributions:**

- Efficient architecture combining depthwise-separable convolutions, squeeze-excitation attention, and residual connections for extreme efficiency (3.82 MB) with minimal degradation (2.04%)
- Few-shot characterization revealing logarithmic scaling, with 500 samples per category achieving 93.5% of full-data performance, providing actionable data collection guidelines
- Augmentation insights showing standard geometric augmentations (+4.52%) perform comparably to elaborate sketch-specific techniques, recommending them as the default
- Domain specialization analysis demonstrating limited cross-dataset transfer (2.09% on ImageNet-Sketch), establishing generalization bounds for efficient models
- Reproducible framework enabling fair comparison and future extensions

References

[1] M. Eitz, J. Hays, and M. Alexa, "How do humans sketch objects?" *ACM Transactions on Graphics*, vol. 31, no. 4, pp. 44:1–44:10, 2012.

[2] R. G. Schneider and T. Tuytelaars, "Sketch classification and classification-driven analysis using fisher vectors," *ACM Transactions on Graphics*, vol. 33, no. 6, pp. 174:1–174:9, 2014.

[3] Q. Yu, Y. Yang, F. Liu, Y.-Z. Song, T. Xiang, and T. M. Hospedales, "Sketch-a-net: A deep neural network that beats humans," *International Journal of Computer Vision*, vol. 122, no. 3, pp. 411–425, 2017.

[4] D. Ha and D. Eck, "A neural representation of sketch drawings," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2018. [Online]. Available: https://openreview.net/forum?id=Hy6GHpkCW

[5] V. Sze, Y.-H. Chen, T.-J. Yang, and J. S. Emer, "Efficient processing of deep neural networks: A tutorial and survey," *Proceedings of the IEEE*, vol. 105, no. 12, pp. 2295–2329, 2017.

[6] D. Ha and D. Eck, "A neural representation of sketch drawings," *arXiv preprint arXiv:1704.03477*, 2017. [Online]. Available: https://arxiv.org/abs/1704.03477

[7] J. Song, K. Pang, Y.-Z. Song, T. Xiang, and T. M. Hospedales, "Learning to sketch with shortcut cycle consistency," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 801–810.

[8] L. Yang, P. Li, X. Qian, H. Su, J. Guo, and J.-J. Zou, "Sketchgnn: Semantic sketch segmentation with graph neural networks," *ACM Transactions on Graphics*, vol. 40, no. 4, pp. 111:1–111:13, 2021.

[9] A. K. Bhunia, P. N. Chowdhury, Y. Yang, T. M. Hospedales, T. Xiang, and Y.-Z. Song, "Vectorization and rasterization: Self-supervised learning for sketch and handwriting," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 5671–5681.

[10] J. Snell, K. Swersky, and R. Zemel, "Prototypical networks for few-shot learning," in *Advances in Neural Information Processing Systems (NIPS)*, vol. 30, 2017, pp. 4077–4087.

[11] C. Finn, P. Abbeel, and S. Levine, "Model-agnostic meta-learning for fast adaptation of deep networks," in *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017, pp. 1126–1135.

[12] K. Pang, K. Li, Y. Yang, H. Zhang, T. M. Hospedales, T. Xiang, and Y.-Z. Song, "Generalising fine-grained sketch-based image retrieval," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 677–686.

[13] Anonymous, "Sketcheffnet: Efficient sketch recognition with deep learning - quickdraw project," January 2025, claude AI Conversation. [Online]. Available: https://claude.ai/share/d8e92f68-b0e1-4629-80d1-30b57ccf762b

[14] N. Dalal and B. Triggs, "Histograms of oriented gradients for human detection," in *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR)*, vol. 1, 2005, pp. 886–893.

[15] M. Eitz, K. Hildebrand, T. Boubekeur, and M. Alexa, "Sketch-based image retrieval: Benchmark and bag-of-features descriptors," *IEEE Transactions on Visualization and Computer Graphics*, vol. 17, no. 11, pp. 1624–1636, 2010.

[16] Y. Cheng, D. Wang, P. Zhou, and T. Zhang, "Model compression and acceleration for deep neural networks: The principles, progress, and challenges," *IEEE Signal Processing Magazine*, vol. 35, no. 1, pp. 126–136, 2018.

[17] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam, "Mobilenets: Efficient convolutional neural networks for mobile vision applications," *arXiv preprint arXiv:1704.04861*, 2017. [Online]. Available: https://arxiv.org/abs/1704.04861

[18] M. Tan and Q. V. Le, "Efficientnet: Rethinking model scaling for convolutional neural networks," in *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019, pp. 6105–6114.

[19] J. Hu, L. Shen, and G. Sun, "Squeeze-and-excitation networks," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 7132–7141.

[20] O. Vinyals, C. Blundell, T. Lillicrap, K. Kavukcuoglu, and D. Wierstra, "Matching networks for one shot learning," in *Advances in Neural Information Processing Systems (NIPS)*, vol. 29, 2016, pp. 3630–3638.

[21] C. Finn, P. Abbeel, and S. Levine, "Model-agnostic meta-learning for fast adaptation of deep networks," in *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017, pp. 1126–1135.

[22] J. Snell, K. Swersky, and R. Zemel, "Prototypical networks for few-shot learning," in *Advances in Neural Information Processing Systems (NIPS)*, vol. 30, 2017, pp. 4077–4087.

[23] K. Weiss, T. M. Khoshgoftaar, and D. Wang, "A survey of transfer learning," *Journal of Big Data*, vol. 3, no. 1, p. 9, 2016.

[24] C. Shorten and T. M. Khoshgoftaar, "A survey on image data augmentation for deep learning," *Journal of Big Data*, vol. 6, no. 1, p. 60, 2019.

[25] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "Imagenet classification with deep convolutional neural networks," in *Advances in Neural Information Processing Systems (NIPS)*, vol. 25, 2012, pp. 1097–1105.

[26] E. D. Cubuk, B. Zoph, D. Mane, V. Vasudevan, and Q. V. Le, "Autoaugment: Learning augmentation strategies from data," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 113–123.

[27] E. D. Cubuk, B. Zoph, J. Shlens, and Q. V. Le, "Randaugment: Practical automated data augmentation with a reduced search space," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2020, pp. 702–703.

[28] P. Y. Simard, D. Steinkraus, and J. C. Platt, "Best practices for convolutional neural networks applied to visual document analysis," in *Proceedings of the Seventh International Conference on Document Analysis and Recognition (ICDAR)*, vol. 2, 2003, pp. 958–963.

[29] H. Wang, S. Ge, Z. Lipton, and E. P. Xing, "Learning robust global representations by penalizing local predictive power," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019, pp. 10 506–10 518.

[30] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778.